

## CSC 555 Final Project Phase 2

Nisarg Patel

### 1) Data Transformation

#### i) Pig

- First copy data from HDFS to Pig Shell as follow:

**grunt>**

lineorder = LOAD 'lineorder.tbl' USING PigStorage('|') AS

(lo\_orderkey:int,lo\_linenummer:int,lo\_custkey:int,lo\_partkey:int,lo\_suppkey:int,lo\_orderdate:int,  
lo\_orderpriority:chararray,lo\_shippriority:chararray,lo\_quantity:int,lo\_extendedprice:int,  
lo\_ordertotalprice:int,lo\_discount:int,lo\_revenue:int,lo\_supplycost:int,lo\_tax:int,  
lo\_commitdate:int,lo\_shipdate:chararray);

- Now store that output file into HDFS Storage with , as delimiter.

**grunt>** store lineorder into '/home/ec2-user/data/lineorder\_csv' using PigStorage(',','-schema');

- Now we will remove the pig schema and then store this file to Local directory

**hadoop fs -rm /home/ec2-user/data/lineorder\_csv/.pig\_schema**

**hadoop fs -getmerge /home/ec2-user/data/lineorder\_csv /home/ec2-user/output.csv**

- Now as you can see from the data head and tail that we got the , (comma) separated file and the columns are no as same as we had in our original file. So we will remove the columns heading that is shown and if we want to use this table we will generate the schema according to this table

```
[ec2-user@ip-172-31-26-136 pig-0.15.0]$ cat output.csv | tail
22657154,1,35206,7987,26259,19980112,4-NOT SPECI,0,26,4926948,11800563,0,4926948,113698,3,19980212,MAIL
22657154,2,35206,225423,6683,19980112,4-NOT SPECI,0,3,404523,11800563,3,392387,80904,0,19980320,REG AIR
22657154,3,35206,281458,33805,19980112,4-NOT SPECI,0,42,6045648,11800563,3,5864278,86366,8,19980329,RAIL
22657155,1,82324,343465,3298,19930414,2-HIGH,0,38,5732110,29573619,8,5273541,90507,5,19930528,AIR
22657155,2,82324,199551,17336,19930414,2-HIGH,0,49,8087695,29573619,4,7764187,99033,5,19930609,AIR
22657155,3,82324,241313,32498,19930414,2-HIGH,0,13,1630590,29573619,2,1597978,75258,2,19930617,RAIL
22657155,4,82324,146210,27369,19930414,2-HIGH,0,11,1381831,29573619,3,1340376,75372,0,19930619,TRUCK
22657155,5,82324,144398,27828,19930414,2-HIGH,0,35,5048365,29573619,1,4997881,86543,2,19930529,SHIP
22657155,6,82324,95079,24535,19930414,2-HIGH,0,47,5048129,29573619,3,4896685,64444,8,19930630,FOB
22657155,7,82324,387303,30452,19930414,2-HIGH,0,19,2641551,29573619,8,2430226,83417,4,19930518,AIR
[ec2-user@ip-172-31-26-136 pig-0.15.0]$ cat output.csv | head
lo_orderkey,lo_linenummer,lo_custkey,lo_partkey,lo_suppkey,lo_orderdate,lo_orderpriority,lo_shippriority,lo_quantity,lo_extendedprice,lo_ordertotalprice,lo_discount,lo_revenue,lo_supplycost,lo_tax,lo_commitdate,lo_shipdate
22657156,1,115205,187139,1097,19930307,3-MEDIUM,0,8,980904,901548,9,892622,73567,1,19930512,TRUCK
22657157,1,35278,229235,37662,19930723,5-LOW,0,42,4889724,4923951,5,4645237,69853,6,19930928,RAIL
22657158,1,104888,340795,24361,19930711,1-URGENT,0,28,5140184,25095892,4,4934576,110146,2,19930823,SHIP
22657158,2,104888,131333,23895,19930711,1-URGENT,0,47,6412351,25095892,10,5771115,81859,2,19930929,RAIL
22657158,3,104888,100012,29888,19930711,1-URGENT,0,33,3339633,25095892,8,3072462,60720,3,19930829,FOB
22657158,4,104888,89669,36648,19930711,1-URGENT,0,28,4644248,25095892,6,4365593,99519,6,19930901,SHIP
22657158,5,104888,252527,14331,19930711,1-URGENT,0,25,3698775,25095892,8,3402873,88770,6,19930819,TRUCK
22657158,6,104888,217181,14917,19930711,1-URGENT,0,20,2196340,25095892,5,2086523,65890,8,19930831,FOB
22657158,7,104888,156623,20613,19930711,1-URGENT,0,3,503886,25095892,2,493808,100777,6,19930928,SHIP
```

- Now we will remove the first line and then print rest of the data.

**awk 'NR>2 {print}' /home/ec2-user/pig-0.15.0/output.csv > /home/ec2-user/pig-0.15.0/lineorder.csv**

ec2-user@ip-172-31-26-136:~/pig-0.15.0

```
[ec2-user@ip-172-31-26-136 pig-0.15.0]$ cat lineorder.csv | head
22657156,1,115205,187139,1097,19930307,3-MEDIUM,0,8,980904,901548,9,892622,73567,1,19930512,TRUCK
22657157,1,35278,229235,37662,19930723,5-LOW,0,42,4889724,4923951,5,4645237,69853,6,19930928,RAIL
22657158,1,104888,340795,24361,19930711,1-URGENT,0,28,5140184,25095892,4,4934576,110146,2,19930823,SHIP
22657158,2,104888,131333,23895,19930711,1-URGENT,0,47,6412351,25095892,10,5771115,81859,2,19930929,RAIL
22657158,3,104888,100012,29888,19930711,1-URGENT,0,33,3339633,25095892,8,3072462,60720,3,19930829,FOB
22657158,4,104888,89669,36648,19930711,1-URGENT,0,28,4644248,25095892,6,4365593,99519,6,19930901,SHIP
22657158,5,104888,252527,14331,19930711,1-URGENT,0,25,3698775,25095892,8,3402873,88770,6,19930819,TRUCK
22657158,6,104888,217181,14917,19930711,1-URGENT,0,20,2196340,25095892,5,2086523,65890,8,19930831,FOB
22657158,7,104888,156623,20613,19930711,1-URGENT,0,3,503886,25095892,2,493808,100777,6,19930928,SHIP
22657159,1,39367,347616,36394,19920119,3-MEDIUM,0,4,665440,800181,5,632168,99816,4,19920229,MAIL
[ec2-user@ip-172-31-26-136 pig-0.15.0]$
```

- File size

```
[ec2-user@ip-172-31-26-136 pig-0.15.0]$ ls -l -h lineorder.csv
-rw-rw-r-- 1 ec2-user ec2-user 2.3G Jun  2 18:53 lineorder.csv
```

## ii) HIVE

- Create table lineorder with the given schema using '|' to separate the column and add the lineorder.tbl file.

```
create table lineorder ( lo_orderkey int, lo_linenummer int, lo_custkey int, lo_partkey int,
lo_suppkey int, lo_orderdate int, lo_orderpriority varchar(15), lo_shippriority varchar(1),
lo_quantity int, lo_extendedprice int, lo_ordertotalprice int, lo_discount int, lo_revenue int,
lo_supplycost int, lo_tax int, lo_commitdate int, lo_shipmode varchar(10) )
```

**ROW FORMAT DELIMITED FIELDS**

**TERMINATED BY ',' STORED AS TEXTFILE;**

- Adding the file to table

**LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl'**

**OVERWRITE INTO TABLE lineorder;**

- Running the query show below.

```
hive> add file /home/ec2-user/addzero.py;
Added resources: [/home/ec2-user/addzero.py]
hive> INSERT OVERWRITE DIRECTORY '/home/ec2-user/data'
> ROW FORMAT DELIMITED
>
> FIELDS TERMINATED BY ','
>
> SELECT TRANSFORM (lo_orderkey, lo_linenummer, lo_custkey, lo_partkey, lo_suppkey, lo_orderdate, lo_orderpriority, lo_shippriority, lo_quantity, lo_extendedprice, lo_ordertotalprice, lo_discount, lo_revenue, lo_supplycost, lo_tax, lo_commitdate, lo_shipmode)
>
> USING 'python addzero.py'
> AS (lo_orderkey, lo_linenummer, lo_custkey, lo_partkey, lo_suppkey, lo_orderdate, lo_orderpriority, lo_shippriority, lo_quantity, lo_extendedprice, lo_ordertotalprice, lo_discount, lo_revenue, lo_supplycost, lo_tax, lo_commitdate, lo_shipmode) FROM lineorder;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. tez, spark) or using Hive 1.X releases.
Query ID = ec2-user_20180607025712_ld4e398a-a006-41dc-a0aa-4af82aad7c9e
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1528309547797_0014, Tracking URL = http://ip-172-31-26-246.us-west-1.compute.internal:8088/proxy/application_1528309547797_0014/
Kill Command = /home/ec2-user/hadoop-2.6.4/bin/hadoop job -kill job_1528309547797_0014
Hadoop job information for Stage-1: number of mappers: 10; number of reducers: 0
2018-06-07 02:57:21,175 Stage-1 map = 0%, reduce = 0%
```

```

Hadoop job information for Stage-1: number of mappers: 10; number of reducers: 0
2018-06-07 03:11:56,852 Stage-1 map = 0%, reduce = 0%
2018-06-07 03:12:57,797 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 27.81 sec
2018-06-07 03:13:58,755 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 85.1 sec
2018-06-07 03:14:59,109 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 139.7 sec
2018-06-07 03:15:51,393 Stage-1 map = 5%, reduce = 0%, Cumulative CPU 188.58 sec
2018-06-07 03:15:54,072 Stage-1 map = 15%, reduce = 0%, Cumulative CPU 189.96 sec
2018-06-07 03:15:55,427 Stage-1 map = 30%, reduce = 0%, Cumulative CPU 191.41 sec
2018-06-07 03:16:56,606 Stage-1 map = 30%, reduce = 0%, Cumulative CPU 247.61 sec
2018-06-07 03:17:57,341 Stage-1 map = 30%, reduce = 0%, Cumulative CPU 303.84 sec
2018-06-07 03:18:49,936 Stage-1 map = 35%, reduce = 0%, Cumulative CPU 353.12 sec
2018-06-07 03:18:52,524 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 355.47 sec
2018-06-07 03:18:53,629 Stage-1 map = 55%, reduce = 0%, Cumulative CPU 355.82 sec
2018-06-07 03:18:54,901 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 356.22 sec
2018-06-07 03:19:32,251 Stage-1 map = 70%, reduce = 0%, Cumulative CPU 369.13 sec
2018-06-07 03:20:33,196 Stage-1 map = 70%, reduce = 0%, Cumulative CPU 398.19 sec
2018-06-07 03:21:33,810 Stage-1 map = 70%, reduce = 0%, Cumulative CPU 426.03 sec
2018-06-07 03:22:30,114 Stage-1 map = 75%, reduce = 0%, Cumulative CPU 452.13 sec
2018-06-07 03:22:31,477 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 452.62 sec
2018-06-07 03:22:32,765 Stage-1 map = 85%, reduce = 0%, Cumulative CPU 453.59 sec
2018-06-07 03:23:32,985 Stage-1 map = 85%, reduce = 0%, Cumulative CPU 480.95 sec
2018-06-07 03:24:33,200 Stage-1 map = 85%, reduce = 0%, Cumulative CPU 509.3 sec
2018-06-07 03:25:26,577 Stage-1 map = 90%, reduce = 0%, Cumulative CPU 534.22 sec
2018-06-07 03:25:28,947 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 534.7 sec
MapReduce Total cumulative CPU time: 8 minutes 54 seconds 700 msec
Ended Job = job_1528309547797_0015
Stage-3 is selected by condition resolver.
Stage-2 is filtered out by condition resolver.
Stage-4 is filtered out by condition resolver.
Moving data to: hdfs://localhost/home/ec2-user/data/.hive-staging_hive_2018-06-07_03-11-48_973_6791712147369275131-1/-ext-10000
Moving data to: /home/ec2-user/data
MapReduce Jobs Launched:
Stage-Stage-1: Map: 10 Cumulative CPU: 534.7 sec HDFS Read: 2417912774 HDFS Write: 2635957096 SUCCESS
Total MapReduce CPU Time Spent: 8 minutes 54 seconds 700 msec
OK
Time taken: 822.15 seconds

```

- This is the generate file to check that go to  
Hadoop fs -ls /home/ec2-user/data

```

[ec2-user@ip-172-31-26-246 ~]$ hadoop fs -ls /home/ec2-user/data
Found 10 items
-rwxr-xr-x 1 ec2-user supergroup 292891221 2018-06-07 03:18 /home/ec2-user/data/000000_0
-rwxr-xr-x 1 ec2-user supergroup 292792102 2018-06-07 03:18 /home/ec2-user/data/000001_0
-rwxr-xr-x 1 ec2-user supergroup 292793330 2018-06-07 03:18 /home/ec2-user/data/000002_0
-rwxr-xr-x 1 ec2-user supergroup 292727002 2018-06-07 03:18 /home/ec2-user/data/000003_0
-rwxr-xr-x 1 ec2-user supergroup 292550431 2018-06-07 03:18 /home/ec2-user/data/000004_0
-rwxr-xr-x 1 ec2-user supergroup 292551004 2018-06-07 03:18 /home/ec2-user/data/000005_0
-rwxr-xr-x 1 ec2-user supergroup 292547728 2018-06-07 03:25 /home/ec2-user/data/000006_0
-rwxr-xr-x 1 ec2-user supergroup 292550746 2018-06-07 03:25 /home/ec2-user/data/000007_0
-rwxr-xr-x 1 ec2-user supergroup 292551063 2018-06-07 03:25 /home/ec2-user/data/000008_0
-rwxr-xr-x 1 ec2-user supergroup 2002469 2018-06-07 03:19 /home/ec2-user/data/000009_0
[ec2-user@ip-172-31-26-246 ~]$ hadoop fs -cat /home/ec2-user/data/000000_0 | head
001,001,29521,310379,16546,19960102,5-LOW,000,017,2361912,18150369,004,2267435,83361,002,19960212,TRUCK
001,002,29521,134619,3259,19960102,5-LOW,000,036,5952996,18150369,009,5417226,99216,006,19960228,MAIL
001,003,29521,127400,1410,19960102,5-LOW,000,009,1141920,18150369,010,1027728,85644,002,19960305,REG AIR
001,004,29521,4263,18842,19960102,5-LOW,000,028,3268328,18150369,009,2974178,70035,006,19960330,AIR
001,005,29521,48054,32491,19960102,5-LOW,000,024,2404920,18150369,010,2164428,60123,004,19960314,FOB
001,006,29521,31269,27344,19960102,5-LOW,000,032,3840832,18150369,007,3571973,72015,002,19960207,MAIL
002,001,62402,212340,21314,19961201,1-URGENT,000,038,4758854,4996796,000,4758854,75139,005,19970114,RAIL
003,001,98653,8594,39169,19931014,5-LOW,000,045,6761655,22702464,006,6355955,90155,000,19940104,AIR
003,002,98653,38071,33331,19931014,5-LOW,000,049,4944443,22702464,010,4449998,60544,000,19931220,RAIL
003,003,98653,256897,28180,19931014,5-LOW,000,027,5005476,22702464,006,4705147,111232,007,19931122,SHIP

```

```

[ec2-user@ip-172-31-26-246 ~]$ hadoop fs -cat /home/ec2-user/data/000009_0 | head
23982182,004,80857,335011,10879,19961209,3-MEDIUM,000,023,2405800,17586382,003,2333626,62760,007,19970118,AIR
23982182,005,80857,129113,29491,19961209,3-MEDIUM,000,006,685266,17586382,004,657855,68526,004,19970120,TRUCK
23982183,001,75700,216120,1197,19951202,1-URGENT,000,026,2693886,8475336,002,2640008,62166,005,19960115,AIR
23982183,002,75700,315456,30492,19951202,1-URGENT,000,019,2795736,8475336,010,2516162,88286,003,19960213,AIR
23982183,003,75700,259318,22241,19951202,1-URGENT,000,021,2682330,8475336,004,2575036,76638,006,19960219,RAIL
23982183,004,75700,348367,309,19951202,1-URGENT,000,003,424605,8475336,010,382144,84921,000,19960120,AIR
23982208,001,43952,354065,39505,19930126,3-MEDIUM,000,001,111905,16835761,004,107428,67143,001,19930307,SHIP
23982208,002,43952,365955,26654,19930126,3-MEDIUM,000,027,5456538,16835761,002,5347407,121256,005,19930225,FOB
23982208,003,43952,193085,6494,19930126,3-MEDIUM,000,033,3887664,16835761,010,3498897,70684,002,19930310,SHIP
23982208,004,43952,366973,15796,19930126,3-MEDIUM,000,036,7343856,16835761,004,7050101,122397,007,19930414,AIR

```

- Python file  
addzero.py

```

GNU nano 2.5.3                               File: addzero.py
./usr/bin/python
import sys

for line in sys.stdin:
    line = line.strip().split('\t')
    lo_orderkey = line[0]
    lo_linenummer = line[1]
    lo_custkey = line[2]
    lo_partkey = line[3]
    lo_supkey = line[4]
    lo_orderdate = line[5]
    lo_orderpriority = line[6]
    lo_shippriority = line[7]
    lo_quantity = line[8]
    lo_extendedprice = line[9]
    lo_ordertotalprice = line[10]
    lo_discount = line[11]
    lo_revenue = line[12]
    lo_supplycost = line[13]
    lo_tax = line[14]
    lo_commitdate = line[15]
    lo_shipmode = line[16]
    try:
        lo_orderkey = '{:03d}'.format(int(lo_orderkey))
        lo_linenummer = '{:03d}'.format(int(lo_linenummer))
        lo_custkey = '{:03d}'.format(int(lo_custkey))
        lo_partkey = '{:03d}'.format(int(lo_partkey))
        lo_supkey = '{:03d}'.format(int(lo_supkey))
        lo_shippriority = '{:03d}'.format(int(lo_shippriority))
        lo_quantity = '{:03d}'.format(int(lo_quantity))
        lo_extendedprice = '{:03d}'.format(int(lo_extendedprice))
        lo_ordertotalprice = '{:03d}'.format(int(lo_ordertotalprice))
        lo_discount = '{:03d}'.format(int(lo_discount))
        lo_revenue = '{:03d}'.format(int(lo_revenue))
        lo_supplycost = '{:03d}'.format(int(lo_supplycost))
        lo_tax = '{:03d}'.format(int(lo_tax))
    except:
        continue

    print '\t'.join([lo_orderkey, lo_linenummer, lo_custkey, lo_partkey, lo_supkey, lo_orderdate, lo_orderpriority, lo_shippriority, lo_quantity, lo_ex

```

### iii) Map Reduce

**hadoop jar ./hadoop-2.6.4/share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input data/addzero/lineorder.tbl -mapper addzero\_mapper.py -file addzero\_mapper.py -reducer addzero\_reducer.py -file addzero\_reducer.py -output data/addzero\_output13**

```

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=2417826195

File Output Format Counters
Bytes Written=2659953700

18/06/09 03:07:24 INFO streaming.StreamJob: Output directory: data/addzero_output13
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -ls data/addzero_output
ls: 'data/addzero_output': No such file or directory
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -ls data/addzero_output13
Found 2 items
-rw-r--r--  2 ec2-user supergroup          0 2018-06-09 03:07 data/addzero_output13/ SUCCESS
-rw-r--r--  2 ec2-user supergroup 2659953700 2018-06-09 03:07 data/addzero_output13/part-000000
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -ls data/addzero_output13/part-000000 | head -5
-rw-r--r--  2 ec2-user supergroup 2659953700 2018-06-09 03:07 data/addzero_output13/part-000000
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -cat data/addzero_output13/part-000000 | head -5
001,003,29521,127400,1418,19960102,5-LOW,000,008,1141920,18150369,010,1027728,85644,002,19960305,REG AIR
001,001,29521,310379,16546,19960102,5-LOW,000,017,2361912,18150369,004,2267435,83361,002,19960212,TRUCK
001,004,29521,4263,18842,19960102,5-LOW,000,028,3268328,18150369,009,2974179,70035,006,19960330,AIR
001,005,29521,48054,32491,19960102,5-LOW,000,024,2404920,18150369,010,2164428,60123,004,19960314,FOB
001,006,29521,31269,27344,19960102,5-LOW,000,032,3840832,18150369,007,3571973,72015,002,19960207,MAIL

```

- Python file for mapper and reducer  
**addzero\_mapper.py**

```

GNU nano 2.5.3                                     File: addzero_mapper.py
#!/usr/bin/python
import sys

for line in sys.stdin:
    line = line.strip().split('|')
    lo_orderkey = line[0]
    lo_linenumber = line[1]
    lo_custkey = line[2]
    lo_partkey = line[3]
    lo_suppkey = line[4]
    lo_orderdate = line[5]
    lo_orderpriority = line[6]
    lo_shippriority = line[7]
    lo_quantity = line[8]
    lo_extendedprice = line[9]
    lo_ordertotalprice = line[10]
    lo_discount = line[11]
    lo_revenue = line[12]
    lo_supplycost = line[13]
    lo_tax = line[14]
    lo_commitdate = line[15]
    lo_shipmode = line[16]
    try:
        lo_orderkey = '{:03d}'.format(int(lo_orderkey))
        lo_linenumber = '{:03d}'.format(int(lo_linenumber))
        lo_custkey = '{:03d}'.format(int(lo_custkey))
        lo_partkey = '{:03d}'.format(int(lo_partkey))
        lo_suppkey = '{:03d}'.format(int(lo_suppkey))
        lo_shippriority = '{:03d}'.format(int(lo_shippriority))
        lo_quantity = '{:03d}'.format(int(lo_quantity))
        lo_extendedprice = '{:03d}'.format(int(lo_extendedprice))
        lo_ordertotalprice = '{:03d}'.format(int(lo_ordertotalprice))
        lo_discount = '{:03d}'.format(int(lo_discount))
        lo_revenue = '{:03d}'.format(int(lo_revenue))
        lo_supplycost = '{:03d}'.format(int(lo_supplycost))
        lo_tax = '{:03d}'.format(int(lo_tax))
    except:
        continue

    print ','.join([lo_orderkey,lo_linenumber,lo_custkey,lo_partkey,lo_suppkey,lo_orderdate,lo_orderpriority,lo_shippriority,lo_quant

```

## addzero\_reducer.py

```

GNU nano 2.5.3                                     File: addzero_reducer.py
#!/usr/bin/python
import sys
import re

for line in sys.stdin:
    print line

```

## 2) Quering

### i) Hive query output

Code :

```

CREATE TABLE lineorder (lo_orderkey int, lo_linenumber int, lo_custkey int, lo_partkey int, lo_suppkey
int, lo_orderdate int, lo_orderpriority varchar(15), lo_shippriority varchar(1), lo_quantity int,
lo_extendedprice int, lo_ordertotalprice int, lo_discount int, lo_revenue int, lo_supplycost int, lo_tax int,
lo_commitdate int, lo_shipmode varchar(10) )
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;

```

```

CREATE TABLE part ( p_partkey int, p_name varchar(22), p_mfgr varchar(6), p_category varchar(7),
p_brand1 varchar(9), p_color varchar(11), p_type varchar(25), p_size int, p_container varchar(10) )

```

ROW FORMAT DELIMITED FIELDS

TERMINATED BY '|' STORED AS TEXTFILE;

```
CREATE TABLE supplier ( s_suppkey int, s_name varchar(25), s_address varchar(25), s_city varchar(10),  
s_nation varchar(15), s_region varchar(12), s_phone varchar(15) )
```

ROW FORMAT DELIMITED FIELDS

TERMINATED BY '|' STORED AS TEXTFILE;

```
LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl'  
OVERWRITE INTO TABLE lineorder;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/part.tbl'  
OVERWRITE INTO TABLE part;
```

```
LOAD DATA LOCAL INPATH '/home/ec2-user/supplier.tbl'  
OVERWRITE INTO TABLE supplier;
```

```
hive> select sum(lo_revenue), p_brand1  
> from lineorder, part, supplier  
> where lo_partkey = p_partkey  
> and lo_suppkey = s_suppkey  
> and p_category = 'MFGR#12'  
> and s_region = 'EUROPE'  
> group by p_brand1;  
WARNING: Hive-on-MR is deprecated in Hive 2  
using Hive 1.X releases.  
Query ID = ec2-user_20180601162308_eedea5e2-  
Total jobs = 2
```

```
Total MapReduce CPU Time Spent: 4 minutes 10 seconds 330 msec  
OK  
11920421610 MFGR#121  
11577268315 MFGR#1210  
13266457869 MFGR#1211  
12181079445 MFGR#1212  
11779690789 MFGR#1213  
11628516109 MFGR#1214  
12721862897 MFGR#1215  
13479058507 MFGR#1216  
13358605308 MFGR#1217  
11219319046 MFGR#1218  
11363358458 MFGR#1219  
12293745533 MFGR#122  
11690136437 MFGR#1220  
11972172012 MFGR#1221  
11804733530 MFGR#1222  
11813329430 MFGR#1223  
13588212841 MFGR#1224  
11997656149 MFGR#1225  
12402823233 MFGR#1226  
11571762246 MFGR#1227  
12722940044 MFGR#1228  
11001916959 MFGR#1229  
13141001739 MFGR#123  
11493744921 MFGR#1230  
11999129834 MFGR#1231  
12883970101 MFGR#1232  
11754484893 MFGR#1233  
11925618535 MFGR#1234  
12228773078 MFGR#1235  
12813848218 MFGR#1236  
12686374495 MFGR#1237  
11013963334 MFGR#1238  
11941917962 MFGR#1239  
12164342558 MFGR#124  
12633409831 MFGR#1240  
11741998486 MFGR#125  
11363542840 MFGR#126  
12633768993 MFGR#127  
12783232889 MFGR#128  
11315290999 MFGR#129  
Time taken: 202.528 seconds, Fetched: 40 row(s)
```

## ii) Pig

```
grunt> supplier = LOAD 'supplier.tbl' USING PigStorage('|') AS (s_suppkey:int,s_name:chararray  
,s_address:chararray,s_city:chararray,s_nation:chararray,s_region:chararray,s_phone:chararray);
```

```
grunt> part = LOAD 'part.tbl' USING PigStorage('|') AS (p_partkey:int, p_name:chararray,  
p_mfgr:chararray,p_category:chararray,p_brand1:chararray,p_color:chararray,p_type:chararray,  
p_size:int,p_container:chararray);
```

```
grunt> lineorder = LOAD 'lineorder.tbl' USING PigStorage('|') AS (lo_orderkey:int,  
lo_linenummer:int,lo_custkey:int,lo_partkey:int,lo_supkey:int,lo_orderdate:int,  
lo_orderpriority:chararray,lo_shippriority:chararray,lo_quantity:int,lo_extendedprice:int,  
lo_ordertotalprice:int,lo_discount:int,lo_revenue:int,lo_supplycost:int,lo_tax:int,  
lo_commitdate:int,lo_shipdate:chararray);
```

```
grunt> group_join = JOIN lineorder BY lo_partkey , part BY p_partkey , supplier BY s_supkey;  
grunt> describe group_join;  
grunt> final_filter = FILTER group_join BY part::p_category == 'MFGR#12' AND supplier::s_region ==  
'EUROPE' ;  
grunt> group_part = group final_filter BY part::p_brand1;  
grunt> group_out = FOREACH group_part GENERATE group ,SUM(final_filter.lo_revenue);
```

```
2018-06-02 23:39:02,555 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!  
2018-06-02 23:39:02,595 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS.  
2018-06-02 23:39:02,596 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.  
2018-06-02 23:39:02,715 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2018-06-02 23:39:02,716 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1  
(MFGR#121,247586631)  
(MFGR#122,874408632)  
(MFGR#123,1408448148)  
(MFGR#124,1794782281)  
(MFGR#125,418846531)  
(MFGR#126,783712452)  
(MFGR#127,358652668)  
(MFGR#128,579786467)  
(MFGR#129,1446006111)  
(MFGR#1210,791012953)  
(MFGR#1211,3342535727)  
(MFGR#1212,1149710128)  
(MFGR#1213,608691518)  
(MFGR#1214,630047208)  
(MFGR#1215,517509064)  
(MFGR#1216,1612803334)  
(MFGR#1217,1111505251)  
(MFGR#1218,1404135378)  
(MFGR#1219,1295295902)  
(MFGR#1220,861583542)  
(MFGR#1221,1514170660)  
(MFGR#1222,1208583671)  
(MFGR#1223,1999016180)  
(MFGR#1224,1965390663)  
(MFGR#1225,661759386)  
(MFGR#1226,1669756821)  
(MFGR#1227,761905381)  
(MFGR#1228,1724881541)  
(MFGR#1229,425834135)  
(MFGR#1230,329053749)  
(MFGR#1231,165356276)  
(MFGR#1232,839635318)  
(MFGR#1233,727370097)  
(MFGR#1234,1499203445)  
(MFGR#1235,853381516)  
(MFGR#1236,567691229)  
(MFGR#1237,1838356790)  
(MFGR#1238,815085135)  
(MFGR#1239,1383003641)  
(MFGR#1240,1125015022)
```

### iii) Map Reduce

This map reduce can be perform in three passes

- i) In first pass we will join lineorder table and part table.
- ii) In second pass we will join the output of first pass and supplier table.
- iii) In third pass atleast add all the revenue that has same brand



First pass

```
hadoop jar ./hadoop-2.6.4/share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input
data/addzero/lineorder.tbl part.tbl -mapper twojoin_mapper1.py -file
twojoin_mapper1.py -reducer twojoin_reducer1.py -file twojoin_reducer1.py -output
data/twojoin_output1
```

```
ec2-user@ip-172-31-26-136:~
Killed map tasks=4
Launched map tasks=23
Launched reduce tasks=1
Data-local map tasks=6
Rack-local map tasks=17
Total time spent by all maps in occupied slots (ms)=9646293
Total time spent by all reduces in occupied slots (ms)=831069
Total time spent by all map tasks (ms)=9646293
Total time spent by all reduce tasks (ms)=831069
Total vcore-milliseconds taken by all map tasks=9646293
Total vcore-milliseconds taken by all reduce tasks=831069
Total megabyte-milliseconds taken by all map tasks=9877804032
Total megabyte-milliseconds taken by all reduce tasks=851014656
Map-Reduce Framework
  Map input records=24596604
  Map output records=24020547
  Map output bytes=559829397
  Map output materialized bytes=607870605
  Input split bytes=2129
  Combine input records=0
  Combine output records=0
  Reduce input groups=408046
  Reduce shuffle bytes=607870605
  Reduce input records=24020547
  Reduce output records=15897
  Spilled Records=59975321
  Shuffled Maps =19
  Failed Shuffles=0
  Merged Map outputs=19
  GC time elapsed (ms)=87757
  CPU time spent (ms)=1732110
  Physical memory (bytes) snapshot=5451751424
  Virtual memory (bytes) snapshot=19789717504
  Total committed heap usage (bytes)=3852468224
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=2468865678
File Output Format Counters
  Bytes Written=510019
18/06/08 22:18:21 INFO streaming.StreamJob: Output directory: data/twojoin_output1
```

```
ec2-user@ip-172-31-26-136:~
GNU nano 2.5.3 File: twojoin_mapper1.py
#!/usr/bin/python
import sys
for line in sys.stdin:
    lo_revenue = '-1'
    lo_suppkey = '-1'
    p_brand = '-1'

    line = line.strip().split(',')
    if len(line) >= 10:
        lo_orderkey = line[0]
        lo_linenum = line[1]
        lo_custkey = line[2]
        partkey = line[3]
        lo_suppkey = line[4]
        lo_revenue = line[12]
        partkey = int(partkey)
        print '%d\t%s\t%s\t%s' % (partkey, p_brand, lo_revenue, lo_suppkey)

    else:
        partkey = line[0]
        p_mfgr = line[3]
        p_brand = line[4]
        partkey = int(partkey)
        if p_mfgr == 'MFGR#12':
            print '%d\t%s\t%s\t%s' % (partkey, p_brand, lo_revenue, lo_suppkey)
```



```
ec2-user@ip-172-31-26-136:~  
GNU nano 2.5.3 File: twojoin_reducer1.py  
#!/usr/bin/python  
import sys  
  
last_key = None  
revenue = ''  
suppkey = ''  
brand = ''  
  
for line in sys.stdin:  
    partkey,p_brand,lo_revenue,p_suppkey = line.strip().split('\t')  
    partkey = int(partkey)  
    if last_key !=partkey:  
        if last_key and suppkey !='-1' and brand!='-1':  
            print '%d\t%s\t%s\t%s\t%s'%(last_key,suppkey,brand,revenue)  
            revenue = lo_revenue  
            last_key = partkey  
            suppkey = p_suppkey  
            brand = p_brand  
    elif last_key == partkey:  
        if p_suppkey !='-1':  
            revenue = lo_revenue  
            suppkey = p_suppkey  
            if p_brand !='-1':  
                brand = p_brand  
        if suppkey !='-1' and brand !='-1':  
            print '%d\t%s\t%s\t%s\t%s'%(last_key,suppkey,brand,revenue)
```

## Second Pass

```
hadoop jar ./hadoop-2.6.4/share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input  
data/twojoin_output1/part-00000 -mapper twojoin_mapper2.py -file  
twojoin_mapper2.py -reducer twojoin_reducer2.py -file twojoin_reducer2.py -output  
data/twojoin_output2
```

```
Map-Reduce Framework  
Map input records=43897  
Map output records=21476  
Map output bytes=689395  
Map output materialized bytes=732365  
Input split bytes=316  
Combine input records=0  
Combine output records=0  
Reduce input groups=21476  
Reduce shuffle bytes=732365  
Reduce input records=21476  
Reduce output records=703  
Spilled Records=42952  
Shuffled Maps =3  
Failed Shuffles=0  
Merged Map outputs=3  
GC time elapsed (ms)=1121  
CPU time spent (ms)=30580  
Physical memory (bytes) snapshot=987824128  
Virtual memory (bytes) snapshot=3972640768  
Total committed heap usage (bytes)=731906048  
  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
  
File Input Format Counters  
Bytes Read=2879469  
File Output Format Counters  
Bytes Written=27405  
18/06/08 22:31:51 INFO streaming.StreamJob: Output directory: data/twojoin_output2  
(ec2-user@ip-172-31-26-136 ~)$
```

```
ec2-user@ip-172-31-26-136:~  
GNU nano 2.5.3 File: twojoin mapper2.py  
#!/usr/bin/python  
import sys  
for line in sys.stdin:  
    partkey = '-1'  
    p_brand = '-1'  
    lo_revenue = '-1'  
    s_region = '-1'  
    line = line.strip().split('\t')  
    if len(line) <= 5:  
        partkey = line[0]  
        suppkey = line[1]  
        p_brand = line[2]  
        lo_revenue = line[3]  
        suppkey = int(suppkey)  
        print '%d\t%s\t%s\t%s\t%s\t' % (suppkey, partkey, p_brand, lo_revenue, s_region)  
    else:  
        suppkey = line[0]  
        region = line[5]  
        suppkey = int(suppkey)  
        if region == 'EUROPE':  
            s_region = 'EUROPE'  
            print '%d\t%s\t%s\t%s\t%s\t' % (suppkey, partkey, p_brand, lo_revenue, s_region)
```

```
ec2-user@ip-172-31-26-136:~  
GNU nano 2.5.3 File: twojoin reducer2.py  
#!/usr/bin/python  
import sys  
last_key = None  
part_key = '-1'  
region = '-1'  
revenue = '-1'  
brand = '-1'  
for line in sys.stdin:  
    line = line.strip().split('\t')  
    suppkey = line[0]  
    partkey = line[1]  
    p_brand = line[2]  
    lo_revenue = line[3]  
    s_region = line[4]  
    suppkey = int(suppkey)  
    if last_key != suppkey:  
        if last_key and part_key != '-1' and brand != '-1' and revenue != '-1' and region != '-1':  
            print '%s\t%d\t%s\t%s\t%s\t' % (last_key, part_key, revenue, brand, region)  
            last_key = suppkey  
            part_key = partkey  
            brand = p_brand  
            revenue = lo_revenue  
            region = s_region  
        elif last_key == suppkey:  
            if region == '-1':  
                region = s_region  
            elif partkey != '-1':  
                part_key = partkey  
                brand = p_brand  
                revenue = lo_revenue  
    if last_key == suppkey:  
        if part_key != '-1' and brand != '-1' and revenue != '-1' and region != '-1':  
            print '%s\t%d\t%s\t%s\t%s\t' % (last_key, part_key, revenue, brand, region)
```

### Third Pass

**hadoop jar ./hadoop-2.6.4/share/hadoop/tools/lib/hadoop-streaming-2.6.4.jar -input data/twojoin\_output2/part-00000 -mapper twojoin\_mapper3.py -file twojoin\_mapper3.py -reducer twojoin\_reducer3.py -file twojoin\_reducer3.py -output data/twojoin\_output3**

```

Map output bytes=12402
Map output materialized bytes=13820
Input split bytes=236
Combine input records=0
Combine output records=0
Reduce input groups=40
Reduce shuffle bytes=13820
Reduce input records=703
Reduce output records=40
Spilled Records=1406
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=820
CPU time spent (ms)=10920
Physical memory (bytes) snapshot=720150528
Virtual memory (bytes) snapshot=2981314560
Total committed heap usage (bytes)=524812288

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=31501
File Output Format Counters
  Bytes Written=751
18/06/08 22:38:21 INFO streaming.StreamJob: Output directory: data/twojoin_output3

```

```

ec2-user@ip-172-31-26-136:~
GNU nano 2.5.3 File: twojoin_mapper3.py

#!/usr/bin/python
import sys
for line in sys.stdin:
    line = line.strip().split('\t')
    last_key = line[0]
    park_key = line[1]
    revenue = line[2]
    brand = line[3]
    print '%s\t%s'%(brand,revenue)

```

```

ec2-user@ip-172-31-26-136:~
GNU nano 2.5.3 File: twojoin_reducer3.py

#!/usr/bin/python
import sys
last_key = ''
total = 0
for line in sys.stdin:
    line = line.strip().split('\t')
    brand = line[0]
    revenue = line[1]
    revenue = int(revenue)
    if last_key !=brand:
        if last_key !='':
            print '%d\t%s'%(total,curr_val)
            last_key = brand
            total = revenue
        elif last_key == brand:
            total +=revenue

    if last_key == brand:
        print '%d\t%s'%(total,last_key)

```

## Output

```

[ec2-user@ip-172-31-26-136 data]$ hadoop fs -cat data/twojoin_output3/part-00000 | head -20
82190141 MFGR#121
72972848 MFGR#1210
81279039 MFGR#1211
31722591 MFGR#1212
76146555 MFGR#1213
34864994 MFGR#1214
84914197 MFGR#1215
87636540 MFGR#1216
89845590 MFGR#1217
75218982 MFGR#1218
48925022 MFGR#1219
60639047 MFGR#122
58936780 MFGR#1220
62603992 MFGR#1221
65066902 MFGR#1222
92845179 MFGR#1223
64458379 MFGR#1224
44999675 MFGR#1225
53217128 MFGR#1226
47684737 MFGR#1227

```

### 3) Clustering

#### A). Using Mahout synthetic clustering

- Now we will generate the sample data using the python notebook as shown below in the figure:

```
In [1]: 1 import pandas as pd
        2 import numpy as np

In [3]: 1 df = pd.DataFrame(np.random.randint(0,100000,size=(10000, 5)), columns=list('ABCDE'))

In [4]: 1 df.shape
Out[4]: (10000, 5)

In [5]: 1 df.head()
Out[5]:
```

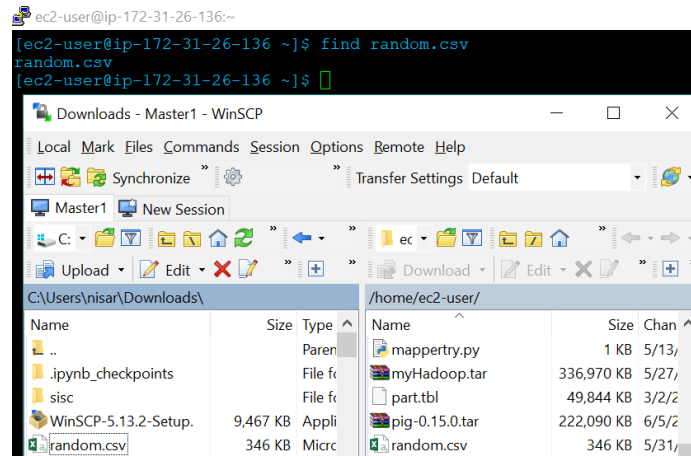
	A	B	C	D	E
0	44054	9540	78383	72681	97176
1	56811	76364	96056	28769	38514
2	80313	42332	64670	17996	41677
3	81432	49217	54326	79925	60445
4	20771	87603	94573	73286	21250

```

In [6]: 1 df.tail()
Out[6]:
```

	A	B	C	D	E
9995	20456	69305	92288	58288	84139
9996	73607	14663	2995	40555	54605
9997	34628	71617	65119	17399	61011
9998	94273	47273	15254	31438	9250
9999	3446	31722	93740	50508	6641

- Now we will generate the output as csv file from python notebook.
  - We will export the csv with ( )space delimiter  
**df.to\_csv("random.csv", sep=' ')**
- Now we will transfer this file from local machine to AWS machine using WinSCP  
Using **find** command we can check if the file is added or not.  
Fig show below:



- Copy this file as text file from csv file .

```
[ec2-user@ip-172-31-26-136 ~]$ cp random.csv random.txt
```

```
[ec2-user@ip-172-31-26-136 ~]$ cat random.txt | head
0 44054 9540 78383 72681 97176
1 56811 76364 96056 28769 38514
2 80313 42332 64670 17996 41677
3 81432 49217 54326 79925 60445
4 20771 87603 94573 73286 21250
5 81295 15680 80926 53568 31304
6 69801 26961 86824 36195 71899
7 45192 92986 65515 91092 23048
8 65905 36871 74303 44642 93844
9 6138 67425 4511 71626 29128
[ec2-user@ip-172-31-26-136 ~]$
```

- This is **another approach** you can also use , delimiter while getting file from python to Aws and then use this code to copy it to new file using space( ) delimiter.

```
~]$ cat random.txt | sed -e s/' '/,/g | cut -d, -f1,2,3,4,5 > random.csv
```

- No we will copy the txt file into the HDFS into testdata folder using command:

```
hadoop fs -put random.txt testdata/
```

```
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -ls testdata
Found 1 items
-rw-r--r-- 2 ec2-user supergroup 353330 2018-06-02 02:35 testdata/random.txt
```

- time mahout org.apache.mahout.clustering.syntheticcontrol.kmeans.Job

Run this code to do K-means iteration's and then note down the time.

```
1.0 : [distance=65241.98512460327]: [9982.0,95983.0,74874.0,90712.0,96891.0,18179.0]
1.0 : [distance=69506.91271608834]: [9999.0,3446.0,31722.0,93740.0,50508.0,6641.0]
18/06/02 03:05:36 INFO ClusterDumper: Wrote 6 clusters
18/06/02 03:05:36 INFO MahoutDriver: Program took 1738058 ms (Minutes: 28.967633333333332)

real    29m48.923s
user    1m25.846s
sys     0m9.211s
[ec2-user@ip-172-31-26-136 ~]$
```

```
[ec2-user@ip-172-31-26-136 ~]$ hadoop fs -ls output
Found 15 items
-rw-r--r-- 2 ec2-user supergroup 194 2018-06-02 03:03 output/_policy
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 03:05 output/clusteredPoints
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:39 output/clusters-0
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:41 output/clusters-1
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 03:03 output/clusters-10-final
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:44 output/clusters-2
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:46 output/clusters-3
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:48 output/clusters-4
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:50 output/clusters-5
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:53 output/clusters-6
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:56 output/clusters-7
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:58 output/clusters-8
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 03:01 output/clusters-9
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:38 output/data
drwxr-xr-x - ec2-user supergroup 0 2018-06-02 02:38 output/random-seeds
[ec2-user@ip-172-31-26-136 ~]$
```

- mahout clusterdump --input output/clusters-10-final --pointsDir output/clusteredPoints --output random\_clusteranalyze.txt

```
[ec2-user@ip-172-31-26-136 ~]$ mahout clusterdump --input output/clusters-10-final --pointsDir output/clusteredPoints --output random_clusteranalyze.txt
Running on hadoop, using /home/ec2-user/hadoop-2.6.4/bin/hadoop and HADOOP_CONF_DIR=
MAHOUT-JOB: /home/ec2-user/apache-mahout-distribution-0.11.2/mahout-examples-0.11.2-job.jar
18/06/02 03:21:31 INFO AbstractJob: Command line arguments: [--dictionaryType=[text], --distanceMeasure=[org.apache.mahout.common.distance.SquaredEuclideanDistanceMeasure], --endPhase=[2147483647], --input=[output/clusters-10-final], --output=[random_clusteranalyze.txt], --outputFormat=[TEXT], --pointsDir=[output/clusteredPoints], --startPhase=[0], --tempDir=[temp]]
18/06/02 03:21:54 INFO ClusterDumper: Wrote 6 clusters
18/06/02 03:21:54 INFO MahoutDriver: Program took 23966 ms (Minutes: 0.3994333333333333)
```

```
[ec2-user@ip-172-31-26-136 ~]$ cat random_clusteranalyze.txt | head
{"r": [2884.599, 25199.675, 22813.472, 17210.01, 20651.808, 23048.789], "c": [4924.985, 65327.102, 63911.634, 75621.71, 31446.28, 64492.532], "n": 1620, "identifier": "CL-265"}
Weight : [props - optional]: Point:
1.0 : [distance=36761.98041557717]: [1.0, 56811.0, 76364.0, 96056.0, 28769.0, 38514.0]
1.0 : [distance=39190.45799796292]: [2.0, 80313.0, 42332.0, 64670.0, 17996.0, 41677.0]
1.0 : [distance=40155.43473939547]: [6.0, 69801.0, 26961.0, 86824.0, 36195.0, 71899.0]
1.0 : [distance=42344.790562129216]: [8.0, 65905.0, 36871.0, 74303.0, 44642.0, 93844.0]
1.0 : [distance=47199.992597234435]: [11.0, 70733.0, 51064.0, 92931.0, 3567.0, 33955.0]
1.0 : [distance=47803.25903222967]: [41.0, 73788.0, 90403.0, 50101.0, 25109.0, 92714.0]
1.0 : [distance=28447.101843770528]: [45.0, 54853.0, 74928.0, 80162.0, 34475.0, 87396.0]
1.0 : [distance=37458.53121005664]: [56.0, 93394.0, 57604.0, 64476.0, 10844.0, 62699.0]
```

## B. Hadoop streaming Map reduce

Creating center.txt file manually with 9 centers

```
GNU nano 2.5.3 File: center.txt
0 1 1
1 10 10
2 100 100
3 500 500
4 1000 1000
5 5000 5000
6 10000 10000
7 50000 50000
8 100000 100000
```

Kmeans\_mapper.py

```
ec2-user@ip-172-31-26-136:~
GNU nano 2.5.3 File: kmeans_mapper.py
#!/usr/bin/python

import sys, re, math

CLUSTERS_FILENAME = 'center.txt'

clusters = []
f = open(clusters_file, 'r')
data = f.read()
f.close()

for line in sys.stdin:
    coord = map(float, line.strip().split())
    distance = (math.sqrt(sum((x-y)**2 for x,y in izip(coord,tab))))
    for tab in table:
        clusters = min(izip(dist , xrange(len(dist))))[1]
    cost = min(dist)**2
    f1 = '%d' % cluster
    f2 = '%4f' % cost
    f3 = ','.join('%4f' % v for v in coord)
    f4 = ','.join((f2,f3))
    print ','.join((f1,f4))
```

